

# Bayesian Data Analysis Assignment 1

Benjamin Cox, S1621312

## Question 1

a)

Our probability vector is  $\theta = (\theta_1, \dots, \theta_6)$  and our outcome vector is  $c = (c_1, \dots, c_6)$ . We are drawing from a multinomial distribution (in the same way that 10 Bern(p) trials are the same as 1 Bin(10,p) trial distributionally), ie

$$c \sim \text{Multinomial}(120, \theta).$$

Therefore the likelihood of  $\theta$  given  $c$  with  $n$  trials is the following:

$$L(\theta|c) = \frac{n!}{c_1!c_2!\dots c_6!} \theta_1^{c_1} \dots \theta_6^{c_6}.$$

A suitable conjugate prior for this would be the Dirichlet distribution (K is the number of possible outcomes, in our case 6),

$$f(x|\alpha, K) = \frac{\Gamma(\sum_{i=1}^K \alpha_i)}{\prod_{i=1}^K \Gamma(\alpha_i)} \prod_{i=1}^K x_i^{\alpha_i-1}.$$

The Jeffrey's prior for the multinomial corresponds to a Dirichlet distribution with

$$\alpha_i = 1/2 \ \forall i \in \{1, \dots, K\}.$$

The above results (conjugacy and Jeffreys prior) are given in [1].

b)

Our posterior distribution for  $\theta$  is

$$\begin{aligned} p(\theta|c) &\propto \left( \frac{\Gamma(3)}{\Gamma(0.5)^6} \theta_1^{-1/2} \dots \theta_6^{-1/2} \right) \left( \frac{n!}{c_1!c_2!\dots c_6!} \theta_1^{c_1} \dots \theta_6^{c_6} \right) \\ &= \text{Dirichlet}(\alpha = c + 0.5, K = 6). \end{aligned}$$

The expected value of the Dirichlet Distribution is given by  $\mathbb{E}[X_i] = \frac{\alpha_i}{\sum \alpha_i}$ , so in our case

$$\mathbb{E}[\theta_i|c] = \frac{c_i + \frac{1}{2}}{\sum c_i + 3}.$$

This corresponds to values of

$\mathbb{E}[\theta_1]$	$\mathbb{E}[\theta_2]$	$\mathbb{E}[\theta_3]$	$\mathbb{E}[\theta_4]$	$\mathbb{E}[\theta_5]$	$\mathbb{E}[\theta_6]$
0.142	0.199	0.183	0.142	0.199	0.134

We are going to compute symmetric 95% credible intervals for each  $\theta_i$ , hence we must marginalise them. We could (theoretically) calculate a 95% credible region in the 6 dimensional parameter space, but this would get extremely complicated really quickly, and would also be hard to interpret.

Fortunately the marginal distributions of the Dirichlet are a lot easier, as they are beta distributions. Write  $\alpha_0 = \sum \alpha_k$ , then we have

$$\theta_i \sim \text{Beta}(\alpha_i, \alpha_0 - \alpha_i).$$

We can substitute in our expressions for  $\alpha_i$  to obtain

$$\theta_i \sim \text{Beta}(c_i + 1/2, c_0 - c_i + 2.5).$$

Using this result we obtain the following 95% credible intervals:

$\theta$	Lower	Upper
$\theta_1$	0.08657456	0.20897331
$\theta_2$	0.1337529	0.2738778
$\theta_3$	0.1200040	0.2556043
$\theta_4$	0.08657456	0.20897331
$\theta_5$	0.1337529	0.2738778
$\theta_6$	0.08007849	0.19945622

c)

We are going to test the null hypothesis of a fair die. For the multinomial distribution there is no easy quantile function, so we are going to use the likelihood-ratio test and Pearson's  $\chi^2$  test. These approach the true  $p$ -value from below and above respectively, so doing both should give us a very good idea.

Computing the  $p$ -value using the likelihood-ratio test we get  $p_{lr} = 0.377$ . Computing the  $p$ -value using Pearson's  $\chi^2$  test we obtain  $p_{\chi^2} = 0.377$ . Thus we can be fairly sure that this is very close to the true  $p$ -value. Therefore we do not reject the null hypothesis that the die is fair.

d)

The posterior predictive distribution is the 'Dirichlet-Multinomial' distribution. The pmf for this is given by

$$f(x|n, \alpha) = \frac{n! \Gamma(\sum \alpha_i)}{\Gamma(n + \sum \alpha_i)} \prod_{k=1}^K \frac{\Gamma(x_k + \alpha_k)}{(x_k!) \Gamma(\alpha_k)}$$

for  $n$  the number of trials and  $\alpha_1, \dots, \alpha_k > 0$ .

Taken as our posterior predictive under the Jeffrey's prior we have

$$c_{\text{new}} \sim \text{DirMNom}(60, c + 1/2).$$

We can simulate from this. We draw 10,000 times from this distribution and find that with probability 0.737 we have more 5s than 6s in our next 60 trials.

e)

We incorporate these into our likelihood, denoting the new count vector as  $d$ . Our new posterior is

$$\theta \sim \text{Dirichlet}(c + d + 1/2, 6).$$

Our new posterior means are

$\mathbb{E}[\theta_1]$	$\mathbb{E}[\theta_2]$	$\mathbb{E}[\theta_3]$	$\mathbb{E}[\theta_4]$	$\mathbb{E}[\theta_5]$	$\mathbb{E}[\theta_6]$
0.163	0.183	0.216	0.138	0.167	0.138

with 95% marginal credible intervals given by

$\theta$	Lower	Upper
$\theta_1$	0.1189814	0.2113785
$\theta_2$	0.1371556	0.2340370
$\theta_3$	0.1667039	0.2698204
$\theta_4$	0.0975285	0.1838316
$\theta_5$	0.1225962	0.2159303
$\theta_6$	0.0939960	0.1791973

The credible intervals have narrowed, as expected for more observations. It is of note that the new credible interval for  $\theta_3$  (barely) does not contain the value required for a ‘fair’ dice. This is evidence that the dice is not fair.

## Question 2

a)

We have two equally credible opinions on the distribution of  $1/\lambda$ . We have one stating that it lies mostly between 5 and 10, and another that states that it lies between 0 and 25. We are going to use the Gamma distribution as our prior for  $\lambda$ .

We are going to calculate parameters for this such that the corresponding inverse gamma distributions has mean of the midpoint and standard deviation of 1/4 the range.

Calculating these parameters we obtain

$$\alpha_1 = 38, \beta_1 = 277.5, \quad \alpha_2 = 6, \beta_2 = 62.5,$$

with the indices corresponding to the originating expert.

This means that our mixture prior is of the form

$$\lambda \sim 0.5 \cdot \Gamma(38, 277.5) + 0.5 \cdot \Gamma(6, 62.5)$$

b)

We have 100 points of data with a mean of 9.877 minutes. We calculate our posterior. The gamma distribution is a conjugate prior to the exponential distribution, so it is quite simple to calculate our new parameters.

We calculate

$$\alpha_1^{post} = 138, \beta_1^{post} = 1265.2, \quad \alpha_2^{post} = 106, \beta_2^{post} = 1050.2$$

using the formulae  $\alpha^{post} = \alpha + n$ ,  $\beta^{post} = \beta + \sum_i x_i = \beta + n\bar{x}$ . To calculate the posterior mixing proportion  $Q$  we use

$$Q = \frac{q \frac{\beta_1^{\alpha_1}}{\Gamma(\alpha_1)} \frac{\Gamma(\alpha_1^{post})}{\beta_1^{post} \alpha_1^{post}}}{q \frac{\beta_1^{\alpha_1}}{\Gamma(\alpha_1)} \frac{\Gamma(\alpha_1^{post})}{\beta_1^{post} \alpha_1^{post}} + (1-q) \frac{\beta_2^{\alpha_2}}{\Gamma(\alpha_2)} \frac{\Gamma(\alpha_2^{post})}{\beta_2^{post} \alpha_2^{post}}}.$$

This quantity involves extremely large numbers, so we work on the log scale and exponentiate at the end.

Overall we obtain a posterior of the form

$$p(\lambda|x) = 0.40455 \cdot \Gamma(138, 1265.2) + 0.59545 \cdot \Gamma(106, 1050.2).$$

Of course the posterior for  $1/\lambda$  is analogous, with inverse gamma rather than gamma.

We calculate a posterior 95% credible interval for  $1/\lambda$  as (8.00, 11.84), and a mean for  $1/\lambda$  of 9.69.

To calculate the posterior probability of waiting for more than 20 minutes we need the posterior predictive distribution. We find that

$$\Pr[x > 20] = 0.127.$$

We note that the posterior distribution is a gamma mixture whereas the posterior predictive distribution is a generalised Pareto distribution. The pdf is given by

$$p(t_*|t) = 0.40455 f_l(\beta_1^{post}, \alpha_1^{post}) + 0.59545 f_l(\beta_2^{post}, \alpha_2^{post}),$$

where

$$f_l(x|a, b) = \frac{a}{b} \left(1 + \frac{x}{b}\right)^{-(a+1)}.$$

This family of distributions is often used in survival analysis to model survival times, so it is natural that it could turn up here.

c)

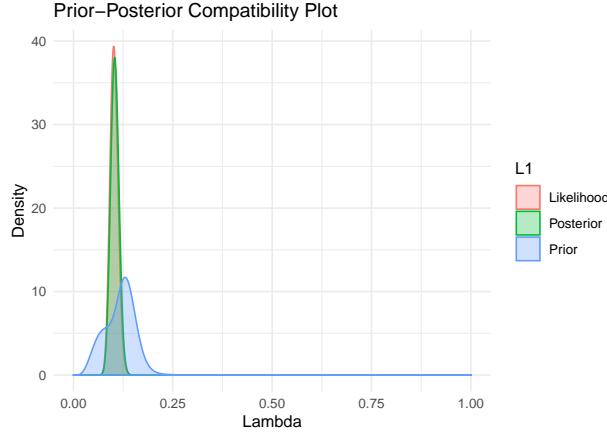


Figure 1: Prior-Likelihood-Posterior Plot

As seen in Figure (1) the posterior is highly influenced by the likelihood and not so much by the prior. The prior is compatible with the data as an initial estimate, and served it's purpose well as the posterior has converged (mostly) onto the likelihood.

The prior is compatible as there is no conflict between the data and the posterior.

d)

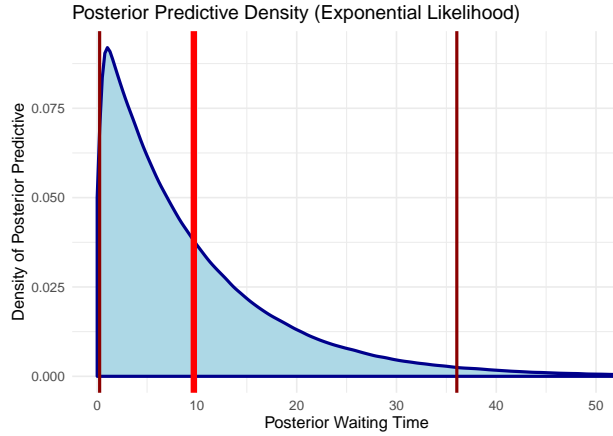


Figure 2: Posterior Predictive Density Plot, bright red line indicates mean, darker red lines indicate 95% CI.

We plot our posterior predictive density. We obtain a mean of 9.69 with symmetric 95% credible interval given by (0.24, 36). This is very much in line with what we expect from the likelihood, indicating good fit.

e)

We are now to take the Lindley distribution as our likelihood. This distribution has pdf of

$$f(t|\lambda) = \frac{\lambda^2}{\lambda + 1} (1 + t) e^{-\lambda t} \implies \ln(f(t|\lambda)) = 2 \ln(\lambda) - \ln(\lambda + 1) + \ln(1 + t) - \lambda t.$$

The log pdf is important for implementing into our MC algorithm.

Given our data we have likelihood of the form

$$L(\lambda|\underline{t}) = \left(\frac{\lambda^2}{\lambda+1}\right)^n \exp\left(-\lambda \sum_{i=1}^n t_i\right) \prod_{i=1}^n (1+t_i),$$

hence we have a posterior of the form

$$\begin{aligned} p(\lambda|\underline{t}) &\propto (\Gamma(38, 277.5) + \Gamma(6, 62.5)) \left(\frac{\lambda^2}{\lambda+1}\right)^n \exp\left(-\lambda \sum_{i=1}^n t_i\right) \prod_{i=1}^n (1+t_i) \\ &\propto \left(\frac{277.5^{38}}{\Gamma(38)} \lambda^{37} e^{-277.5 \cdot \lambda} + \frac{62.5^6}{\Gamma(6)} \lambda^5 e^{-62.5 \cdot \lambda}\right) \left(\frac{\lambda^2}{\lambda+1}\right)^n \exp\left(-\lambda \sum_{i=1}^n t_i\right) \prod_{i=1}^n (1+t_i) \end{aligned}$$

We calculate our posterior. We obtain an expected waiting time of 10.4 minutes, an increase from our previous.

We also plot the new posterior predictive.

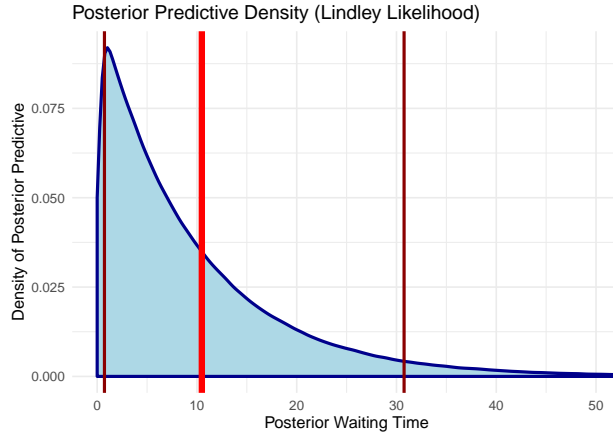


Figure 3: Posterior Predictive Density Plot for Lindley Likelihood, bright red line indicates mean, darker red lines indicate 95% CI.

Notice that the credible interval for the posterior predictive derived from the Lindley likelihood are a lot tighter.

f)

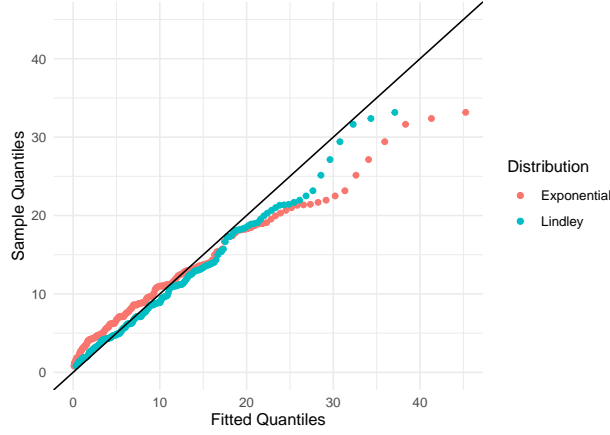


Figure 4: Q-Q plot of sample quantiles against posterior predictive quantiles

We see that the Lindley distribution fits the data much better than the exponential distribution. We perform Kolmogorov-Smirnoff tests comparing our posterior predictives and the data. For the Lindley likelihood we obtain a distance of 0.0685 with  $p$ -value of 0.736. For the exponential likelihood we obtain a distance of 0.1797 with corresponding  $p$ -value of 0.003.

This is extremely good evidence that the Lindley distribution is a better fit for the data than the Exponential distribution.

### Question 3

a)

In order to make our results directly comparable we are going to mean centre the data before we fit the initial linear model. This should only affect the intercept term.

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.2439	0.0643	3.79	0.0002
SexI	-0.8249	0.1024	-8.06	0.0000
SexM	0.0577	0.0833	0.69	0.4887
Length	-0.4583	1.8091	-0.25	0.8000
Diameter	11.0751	2.2273	4.97	0.0000
Height	10.7615	1.5362	7.01	0.0000
Whole.weight	8.9754	0.7254	12.37	0.0000
Shucked.weight	-19.7869	0.8174	-24.21	0.0000
Viscera.weight	-10.5818	1.2937	-8.18	0.0000
Shell.weight	8.7418	1.1247	7.77	0.0000

This points to length and Male-Female difference not having a significant effect on age. Length is likely due to high correlation with height (Pearson correlation of 0.828) and diameter (Pearson correlation of 0.987). Most of the variables have extremely high correlation, so this is going to require a long MCMC run.

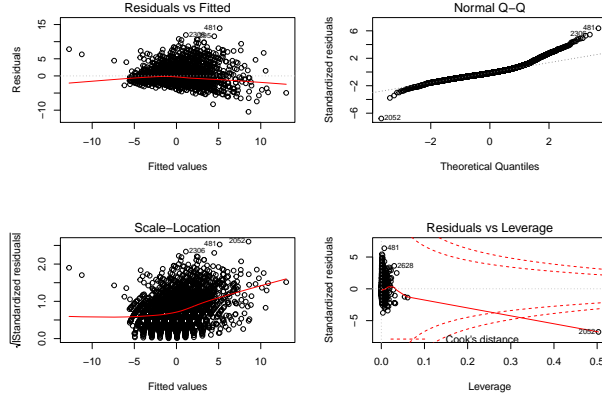


Figure 5: Diagnostic plots for the standard linear model

Looking at diagnostic plots we see a large number of extreme observations with high leverage. Therefore it makes sense to perform a robust regression using  $t$  distributed errors.

b)

We are going to fit a Bayesian linear model of the form

$$\begin{aligned}
 y &\sim t_{\nu+1}(X\beta, \sigma), \\
 \beta_i &\sim \text{Normal}(0, 10^4), \\
 \sigma &\sim \Gamma^{-1}(0.1, 0.1), \\
 \nu &\sim \text{Exponential}(0.034),
 \end{aligned}$$

where  $X$  is the  $n \times p$  centred design matrix,  $\beta$  is the  $p$  vector of regression coefficients, and the  $t$  distribution is understood to be shifted and scaled such that  $\frac{y - X\beta}{\sigma}$  has a classic  $t$  distribution with  $\nu + 1$  degrees of freedom.

c)

Fitting the model we obtain the following quantities:

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
beta[1]	-0.22	0.00	0.06	-0.33	-0.25	-0.22	-0.18	-0.11	28239.55	1.00
beta[2]	-0.67	0.00	0.08	-0.83	-0.73	-0.67	-0.61	-0.51	32862.97	1.00
beta[3]	0.12	0.00	0.07	-0.02	0.07	0.12	0.17	0.26	33774.18	1.00
beta[4]	1.88	0.01	1.51	-1.13	0.86	1.89	2.90	4.82	31346.96	1.00
beta[5]	7.28	0.01	1.86	3.64	6.02	7.27	8.53	10.96	31154.47	1.00
beta[6]	15.80	0.01	2.04	11.81	14.43	15.80	17.17	19.80	50288.95	1.00
beta[7]	7.51	0.01	0.80	5.94	6.97	7.51	8.04	9.06	23090.61	1.00
beta[8]	-16.35	0.01	0.90	-18.12	-16.95	-16.35	-15.74	-14.59	26678.37	1.00
beta[9]	-9.22	0.01	1.20	-11.56	-10.03	-9.22	-8.41	-6.89	34525.85	1.00
beta[10]	6.72	0.01	1.17	4.43	5.92	6.71	7.50	9.05	28005.24	1.00
sigma	1.39	0.00	0.03	1.33	1.37	1.39	1.41	1.45	35667.27	1.00
nu	1.87	0.00	0.16	1.58	1.76	1.86	1.97	2.19	35705.46	1.00
y_new[1]	-0.34	0.01	2.47	-4.85	-1.41	-0.34	0.74	4.22	56374.07	1.00
lp_	-6470.43	0.02	2.47	-6476.21	-6471.85	-6470.08	-6468.65	-6466.64	21223.39	1.00

Table 1: Regression Coefficient Table

d)

The model table contains our  $\hat{R}$  values, and shows that we have good reason to believe that the chains have converged.

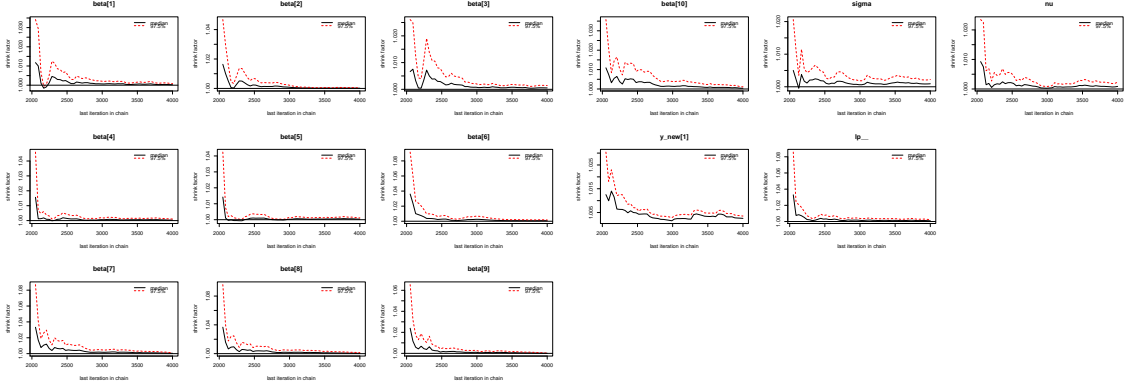


Figure 6: Plots of  $\hat{R}$  as chains progressed

Although Figure 6 is a bit small it is clear that after our burn-in  $\hat{R}$  was very close to 1, and got closer as we computed more iterations.

e)

We are going to alter some of the prior hyperparameters in order to test sensitivity. We are going to assume the following model:

$$\begin{aligned} y &\sim t_\nu(X\beta, \sigma), \\ \beta_i &\sim \text{Normal}(0, 10), \\ \sigma &\sim \Gamma^{-1}(7, 30), \\ \nu &\sim \text{Exponential}(0.034), \end{aligned}$$

The hyperparameters are quite different, so this should be a good test of prior sensitivity. We are going to run the model for less iterations as we do not intend to use this model for inference, merely for prior sensitivity analysis.

We obtain the following table of results:

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
beta[1]	-0.21	0.00	0.06	-0.32	-0.25	-0.21	-0.17	-0.10	5572.99	1.00
beta[2]	-0.68	0.00	0.08	-0.84	-0.73	-0.68	-0.62	-0.51	6239.80	1.00
beta[3]	0.12	0.00	0.07	-0.02	0.07	0.12	0.17	0.25	6566.89	1.00
beta[4]	1.98	0.02	1.47	-0.89	0.97	1.98	2.99	4.82	6567.40	1.00
beta[5]	7.23	0.02	1.80	3.71	5.98	7.24	8.46	10.72	6482.90	1.00
beta[6]	15.25	0.02	2.01	11.29	13.91	15.27	16.63	19.13	9565.22	1.00
beta[7]	7.36	0.01	0.78	5.82	6.83	7.36	7.89	8.86	4357.55	1.00
beta[8]	-16.20	0.01	0.90	-17.95	-16.81	-16.20	-15.60	-14.44	4957.90	1.00
beta[9]	-8.98	0.01	1.19	-11.31	-9.78	-9.00	-8.19	-6.64	6498.79	1.00
beta[10]	6.93	0.02	1.15	4.74	6.15	6.91	7.70	9.25	5590.50	1.00
sigma	1.40	0.00	0.03	1.34	1.38	1.40	1.42	1.46	6930.44	1.00
nu	1.90	0.00	0.16	1.61	1.79	1.89	2.00	2.22	7134.90	1.00
y_new[1]	-0.36	0.02	2.54	-4.77	-1.41	-0.35	0.72	4.13	10775.23	1.00
lp--	-6497.93	0.04	2.39	-6503.44	-6499.35	-6497.65	-6496.19	-6494.16	4370.59	1.00

This is nearly identical to the table obtained previously, so we can be confident in our results. All of the differences in means are well within their standard deviations, so overall we are happy.



f)

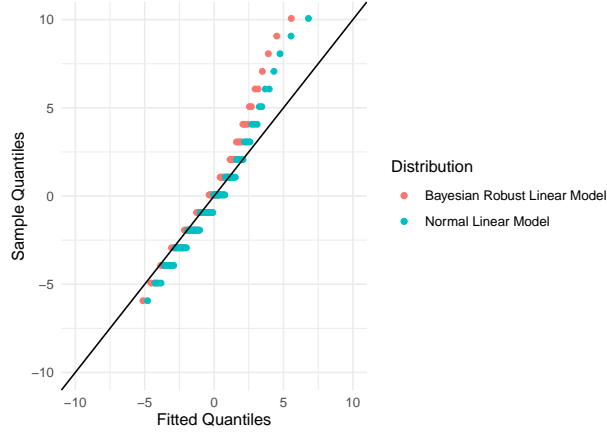


Figure 7: QQ plots comparing our models

Looking at figure 7 neither of our models fit the data very well. Performing KS tests we find a distance of 0.11827 between our predictions from the Robust Bayesian and the true data, compared with a distance of 0.15466 for the normal linear and the true data. We note that the true data is rather skewed with estimated skewness 1.11. This has not been compensated for in the model. This means that there are more old abalones than predicted by the models. This is likely due to their characteristics not changing that much after they age enough. This is shown by the deviation of the upper quantiles of the predictions from the sample quantiles.

g)

95% credible intervals are given in Table 1. We see that the credible intervals for  $\beta_3$  and  $\beta_4$  contain 0. These are the coefficients linked to 'isMale' and Length respectively. This observation is identical to that of the normal linear model.

h)

The 95% credible interval for the age relating to the new data is given in the Table 1 as  $(-4.85, 4.22)$  with mean of  $-0.34$ . We plot the density of the posterior predictive and obtain the following:

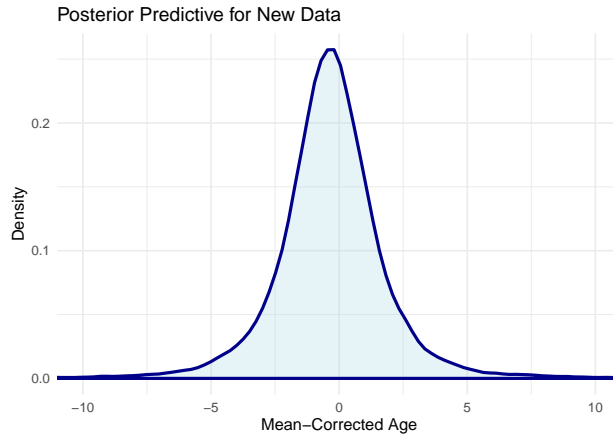


Figure 8: Posterior predictive density for new data

## References

- [1] Frank Tuyl. “A Note on Priors for the Multinomial Model”. In: *The American Statistician* 71.4 (2017), pp. 298–301.

## A R Code for Question 1

```
1 library(extraDistr)
2
3 obs <- c(17, 24, 22, 17, 24, 16)
4 obs2 <- c(22, 20, 30, 16, 16, 16)
5 value <- 1:6
6 names(obs) <- c("1", "2", "3", "4", "5", "6")
7 names(obs2) <- names(obs)
8 n1 <- sum(obs)
9 n2 <- sum(obs2)
10 ysum <- sum(obs * value)
11 #categorical
12 #jeffreys prior is dirichlet(a = 1/2)
13
14 #theta_1, ..., theta_6 ~ Dirichlet(a_1, ..., a_6)
15 #c_1, ..., c_6 ~ Multinomial(theta_1, ..., theta_6)
16 #a'_k = a_k + ysum
17
18 j_prior_params <- rep(1 / 2, 6)
19
20 post_params <- j_prior_params + obs
21 post_params_new <- j_prior_params + obs + obs2
22
23 post_pred <- post_params / sum(post_params)
24 post_pred_new <- post_params_new / sum(post_params_new)
25
26 alpha = 0.05
27 for (i in post_params){
28   print(qbeta(c(0.025, 0.975), i, sum(post_params) - i))
29 }
30 for (i in post_params_new){
31   print(qbeta(c(0.025, 0.975), i, sum(post_params_new) - i))
32 }
33
34 # n <- 1e6
35 # rope_r <- 0.05*1/30
36 # A <- rdirichlet(n, post_params)
37 # B <- as.matrix(A)
38 # f <- function(x) all((x > 1/6-rope_r)*(x < 1/6+rope_r))
39 # mean(apply(B, MARGIN = 1, FUN = f))
40
41 # nh_p <- rep(1/6, 6)
42 # n <- 1e2
43 # size <- sum(obs)
44 # nh_lik <- t(rmultinom(n, size, nh_p))
45
46 #pearson chi-squared
47
48 expected <- n1 * rep(1/6, 6)
49 observed <- obs
50 sq_diff <- (observed - expected)^2
51 pear_stat <- sq_diff / expected
52 chi_stat <- sum(pear_stat)
53 chi_dof <- length(obs) - 1
54 chi_p <- pchisq(chi_stat, chi_dof)
55 chi_p
56
57 nh <- rep(1/6, 6)
58 mle <- obs/n1
59 lr_stat <- -2 * sum(obs * log(nh/mle))
60 lr_chi_stat <- pchisq(lr_stat, chi_dof)
61 lr_chi_stat
62
63 n <- 1e5
64
65 C <- rdirnmom(n, 60, post_params)
66 D <- as.matrix(C)
67 ts <- D[,5] > D[,6]
68 sum(ts)/n
```

## B R Code for Question 2

```
1 library(tidyverse)
2 library(HDIInterval)
3 library(extraDistr)
4
5 par(mfrow = c(1, 1))
6
7 waiting_times <-
8   c(
9     0.8,
10    0.8,
11    1.3,
12    1.5,
13    1.8,
14    1.9,
```

```

15 | 1.9,
16 | 2.1,
17 | 2.6,
18 | 2.7,
19 | 2.9,
20 | 3.1,
21 | 3.2,
22 | 3.3,
23 | 3.5,
24 | 3.6,
25 | 4.0,
26 | 4.1,
27 | 4.2,
28 | 4.2,
29 | 4.3,
30 | 4.3,
31 | 4.4,
32 | 4.4,
33 | 4.6,
34 | 4.7,
35 | 4.7,
36 | 4.8,
37 | 4.9,
38 | 4.9,
39 | 5,
40 | 5.3,
41 | 5.5,
42 | 5.7,
43 | 5.7,
44 | 6.1,
45 | 6.2,
46 | 6.2,
47 | 6.2,
48 | 6.3,
49 | 6.7,
50 | 6.9,
51 | 7.1,
52 | 7.1,
53 | 7.1,
54 | 7.1,
55 | 7.4,
56 | 7.6,
57 | 7.7,
58 | 8,
59 | 8.2,
60 | 8.6,
61 | 8.6,
62 | 8.6,
63 | 8.8,
64 | 8.8,
65 | 8.9,
66 | 8.9,
67 | 9.5,
68 | 9.6,
69 | 9.7,
70 | 9.8,
71 | 10.7,
72 | 10.9,
73 | 11,
74 | 11,
75 | 11.1,
76 | 11.2,
77 | 11.2,
78 | 11.5,
79 | 11.9,
80 | 12.4,
81 | 12.5,
82 | 12.9,
83 | 13,
84 | 13.1,
85 | 13.3,
86 | 13.6,
87 | 13.7,
88 | 13.9,
89 | 14.1,
90 | 15.4,
91 | 15.4,
92 | 17.3,
93 | 17.3,
94 | 18.1,
95 | 18.2,
96 | 18.4,
97 | 18.9,
98 | 19,
99 | 19.9,
100 | 20.6,
101 | 21.3,
102 | 21.4,
103 | 21.9,
104 | 23.0,
105 | 27,
106 | 31.6,
107 | 33.1,
108 | 38.5
109 | )
110 |
111 | gammaShRaFromMeanSD = function(mean , sd) {
112 |   if (mean <= 0)
113 |     stop("'mean must be > 0'")
114 |   if (sd <= 0)
115 |     stop("'sd must be > 0'")
116 |   shape = mean ^ 2 / sd ^ 2
117 |   rate = mean / sd ^ 2
118 |   return(list(shape = shape , rate = rate))
119 | }
120 |
121 | gammaShRaFromModeSD = function(mode , sd) {

```

```

122 if (mode <= 0)
123   stop("mode must be > 0")
124 if (sd <= 0)
125   stop("sd must be > 0")
126 rate = (mode + sqrt(mode ^ 2 + 4 * sd ^ 2)) / (2 * sd ^ 2)
127 shape = 1 + mode * rate
128 return(list(shape = shape, rate = rate))
129 }
130
131 iGammaShScFromMeanVar = function(mean, var) {
132   shape <- (mean ^ 2 / var) + 2
133   scale <- (mean ^ 3 / var) + mean
134   return(list(shape = shape, scale = scale))
135 }
136
137 ex_1_prior <- iGammaShScFromMeanVar(7.5, (0.25 * (10 - 5)) ^ 2)
138 ex_2_prior <- iGammaShScFromMeanVar(12.5, (0.25 * (25 - 0)) ^ 2)
139 #prior is 0.5G(9, 0.6) + 0.5G(1,0.08)
140
141 n <- length(waiting_times)
142 mean_wait <- mean(waiting_times)
143
144 #posterior is Q*G(9+n, 0.6+n*xbar) + (1-Q)*G(12.5+n, 12.5+n*xbar)
145
146 p <- 0.5
147
148 find_mix_coefs_2gamma <- function(a1, b1, a2, b2, q1, n, xbar) {
149   q2 <- 1 - q1
150   a1post <- a1 + n
151   a2post <- a2 + n
152   b1post <- b1 + (n * xbar)
153   b2post <- b2 + (n * xbar)
154   log_c1 <-
155     (a1 * log(b1) - lgamma(a1)) + (lgamma(a1post) - (a1post * log(b1post)))
156   log_c2 <-
157     (a2 * log(b2) - lgamma(a2)) + (lgamma(a2post) - (a2post * log(b2post)))
158   log_Q <-
159     (log(q1) + log_c1) - (log(q1) + log_c1 + log(1 + q2 / q1 * exp(log_c2 -
160       log_c1)))
161   #log_Q <- log(q1) + log_c1 - log(q1 * exp(log_c1) + q2 * exp(log_c2))
162   Q <- exp(log_Q)
163   # c1 <- ((b1^a1)/(gamma(a1))) * ((gamma(a1post))/(b1post^a1post))
164   # c2 <- ((b2^a2)/(gamma(a2))) * ((gamma(a2post))/(b2post^a2post))
165   # Q <- (q1 * c1)/(q1*c1 + q2*c2)
166   return(c(Q, 1 - Q))
167 }
168
169 mcoefs <-
170   find_mix_coefs_2gamma(ex_1_prior[[1]],
171     ex_1_prior[[2]],
172     ex_2_prior[[1]],
173     ex_2_prior[[2]],
174     p,
175     n,
176     mean_wait)
177
178 lambda <- seq(0, 1, len = 5000)
179
180 likelihood <-
181   as.matrix(apply(as.array(waiting_times), 1, dexp, lambda))
182 likelihood <- apply(likelihood, 1, prod)
183
184 area = sfsmisc::integrate.xy(lambda, likelihood)
185 const = 1 / area
186 likelihood <- const * likelihood
187
188 prior <-
189   p * dgamma(lambda, ex_1_prior[[1]], ex_1_prior[[2]]) + (1 - p) * dgamma(lambda, ex_2_prior[[1]], ex_2_prior[[2]])
190 posterior <-
191   mcoefs[[1]] * dgamma(lambda, ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
192     mean_wait) + mcoefs[[2]] * dgamma(lambda, ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
193     mean_wait)
194
195 plot(
196   lambda,
197   prior,
198   col = "darkgreen",
199   ylab = "Density",
200   xlab = expression(lambda),
201   type = "l",
202   ylim = c(0, 45),
203   lwd = 2
204 )
205 lines(lambda, likelihood, col = "blue2", lwd = 2)
206 lines(lambda, posterior, col = "red", lwd = 2)
207 legend(
208   "topright",
209   legend = c("prior", "scaled likelihood", "posterior"),
210   lty = c(1, 1, 1),
211   lwd = c(2, 2, 2),
212   col = c("darkgreen", "blue2", "red"),
213   bty = "n"
214 )
215
216 #prior is relatively flat, posterior is very compatible with data and not very influenced by prior
217
218 post_mean <-
219   mcoefs[[1]] * ((ex_1_prior[[1]] + n) / (ex_1_prior[[2]] + n * mean_wait)) + mcoefs[[2]] *
220   ((ex_2_prior[[1]] + n) / (ex_2_prior[[2]] + n * mean_wait))
221 post_mean_wait <- 1 / post_mean
222
223 posterior_ci <-
224   mcoefs[[1]] * qgamma(c(0.025, 0.975), ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
225     mean_wait) + mcoefs[[2]] * qgamma(c(0.025, 0.975), ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
226     mean_wait)
227
228 pci_area <-

```

```

229 sfsmisc::integrate.xy(lambda, posterior, posterior_ci[[1]], posterior_ci[[2]])
230
231
232 time <- seq(0, 30, len = 5000)
233
234 post_wait <-
235   mcoefs[[1]] * dinvgamma(time, ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
236     mean_wait) + mcoefs[[2]] * dinvgamma(time, ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
237     mean_wait)
238
239 post_wait_ci <-
240   mcoefs[[1]] * qinvgamma(c(0.025, 0.975), ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
241     mean_wait) + mcoefs[[2]] * qinvgamma(c(0.025, 0.975), ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
242     mean_wait)
243
244 post_wait_20 <-
245   mcoefs[[1]] * pinvgamma(20, ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
246     mean_wait) + mcoefs[[2]] * pinvgamma(20, ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
247     mean_wait)
248
249 post_wait_sim <-
250   mcoefs[[1]] * rinvgamma(1e5, ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
251     mean_wait) + mcoefs[[2]] * rinvgamma(1e5, ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
252     mean_wait)
253
254 N <- 1e6
255
256 components <-
257   sample(
258     1:2,
259     prob = c(mcoefs[[1]], mcoefs[[2]]),
260     size = N,
261     replace = TRUE
262   )
263
264 alphas <- c(ex_1_prior[[1]] + n, ex_2_prior[[1]] + n)
265 betas <-
266   c(ex_1_prior[[2]] + n * mean_wait, ex_2_prior[[2]] + n * mean_wait)
267
268 samples <- rinvgamma(N, alphas[components], betas[components])
269
270 plot(density(samples))
271 mean(samples)
272 quantile(samples, c(0.025, 0.975))
273
274 post_pred <-
275   Renext::rlomax(N, betas[components], alphas[components])
276
277 plot(density(post_pred))
278 mean(post_pred)
279 quantile(post_pred, c(0.025, 0.975))
280 sum(post_pred > 20) / N
281
282 library(coda)
283 library(rstan)
284
285 rstan_options(auto_write = TRUE)
286
287 b_mm_data <-
288   list(
289     K = 2,
290     N = n,
291     y = waiting_times,
292     theta = c(0.5, 0.5),
293     alpha = c(6, 38),
294     beta = c(62.5, 277.5)
295   )
296
297 # b_data <- list(N=n, y = waiting_times, gprior=c(38,278))
298
299 # b_m <- stan_model(file = 'aiq2simp.stan')
300
301 # b_m_fit <- sampling(
302 #   b_m,
303 #   data = b_data,
304 #   chains = 7,
305 #   control = list(adapt_delta = 0.8),
306 #   iter = 1e5
307 # )
308 # plot(b_m_fit)
309 # b_m_fit
310
311 # smps <- extract(b_m_fit)
312 # tp <- traceplot(b_m_fit, pars = c("lambda"))
313
314 # b_m_coda <- As.mcmc.list(b_m_fit)
315 # gelman.plot(b_m_coda, ask=FALSE)
316 # gelman.diag(b_m_coda)
317 # plot(density(smps$postdraw))
318 # sum(smps$postdraw > 20) / (length(smps$postdraw))
319 # mean(smps$postdraw)
320 # quantile(smps$postdraw, c(0.025, 0.975))
321 # plot(density(smps$ewt))
322 # sum(smps$ewt > 20) / (length(smps$ewt))
323 # mean(smps$ewt)
324 # quantile(smps$ewt, c(0.025, 0.975))
325
326 b_mm <- stan_model(file = 'aiq2.stan')
327
328 b_mm_fit <- sampling(
329   b_mm,
330   data = b_mm_data,
331   chains = 7,
332   control = list(adapt_delta = 0.8),
333   iter = 40000
334 )
335 plot(b_mm_fit)
336 b_mm_fit
337

```

```

336 smps <- extract(b_mm_fit)
337 tp <- traceplot(b_mm_fit, pars = c("lambda", "mwt"))
338 tp
339
340 b_mm_coda <- As.mcmc.list(b_mm_fit)
341 gelman.plot(b_mm_coda, ask = FALSE)
342 gelman.diag(b_mm_coda)
343
344 plot(density(smps$postdraw))
345 sum(smps$postdraw > 20) / (length(smps$postdraw))
346 mean(smps$postdraw)
347 quantile(smps$postdraw, c(0.025, 0.975))
348
349 plot(density(smps$mwt))
350 sum(smps$mwt > 20) / (length(smps$mwt))
351 mean(smps$mwt)
352 quantile(smps$mwt, c(0.025, 0.975))
353
354 require(rjags)
355
356 model = jags.model(
357   file = "aiq2jags.jags",
358   data =
359     list(
360       y = waiting_times,
361       n = n,
362       a = c(38, 6),
363       b = c(277.5, 62.5),
364       p = c(0.5, 0.5)
365     ),
366   n.chains = 10
367 )
368
369 # Burnin for 1000 samples
370 update(model, 100000, progress.bar = "none")
371
372 # Running the model
373 res = coda.samples(
374   model,
375   variable.names = c("lambda", "ypred", "texp"),
376   n.iter = 200000,
377   progress.bar = "none"
378 )
379 summary(res)
380
381 qqplot(smps$postdraw, waiting_times, xlim = c(0, 40))
382 qqplot(post_pred, waiting_times, xlim = c(0, 40))
383 qqplot(res[[1]][, "ypred"], waiting_times, xlim = c(0, 40))
384
385
386 b_mm_lind <- stan_model(file = 'aiq2lindley.stan')
387
388 b_mm_lind_fit <- sampling(
389   b_mm_lind,
390   data = b_mm_data,
391   chains = 7,
392   control = list(adapt_delta = 0.8),
393   iter = 40000
394 )
395 plot(b_mm_lind_fit)
396 b_mm_lind_fit
397
398 smps_lind <- extract(b_mm_lind_fit)
399 tp_lind <- traceplot(b_mm_lind_fit, pars = c("lambda", "mwt"))
400 tp_lind
401
402 b_mm_lind_coda <- As.mcmc.list(b_mm_lind_fit)
403 gelman.plot(b_mm_lind_coda, ask = FALSE)
404 gelman.diag(b_mm_lind_coda)
405
406 plot(density(smps_lind$postdraw))
407 sum(smps_lind$postdraw > 20) / (length(smps_lind$postdraw))
408 mean(smps_lind$postdraw)
409 quantile(smps_lind$postdraw, c(0.025, 0.975))
410
411 plot(density(smps_lind$mwt))
412 sum(smps_lind$mwt > 20) / (length(smps_lind$mwt))
413 mean(smps_lind$mwt)
414 quantile(smps_lind$mwt, c(0.025, 0.975))
415
416 qqplot(smps_lind$postdraw, waiting_times, xlim = c(0, 40))
417 qqplot(smps$postdraw, waiting_times, xlim = c(0, 40))
418 ks.test(smps_lind$postdraw, waiting_times)
419 ks.test(smps$postdraw, waiting_times)
420
421
422 dens <- data.frame(postdraw = post_pred)
423 ggdens <-
424   ggplot(data = dens, aes(x = postdraw)) +
425     geom_density(size = 1, color = "darkblue", fill = "lightblue") +
426     theme_minimal() +
427     labs(x = "Posterior Waiting Time", y = "Density of Posterior Predictive", title = "Posterior Predictive Density (Exponential Likelihood)") +
428     geom_vline(
429       xintercept = mean(dens$postdraw),
430       size = 2,
431       color = "red"
432     ) +
433     geom_vline(
434       xintercept = quantile(dens$postdraw, c(0.025, 0.975)),
435       color = "darkred",
436       size = 1
437     ) + coord_cartesian(xlim = c(0, 50))
438 ggdens
439
440 lambda <- seq(0, 1, len = 5000)
441

```

```

443 likelihood <-
444   as.matrix(apply(as.array(waiting_times), 1, dexp, lambda))
445 likelihood <- apply(likelihood, 1, prod)
446
447 area = sfsmisc::integrate.xy(lambda, likelihood)
448 const = 1 / area
449 likelihood <- const * likelihood
450
451 prior <-
452   p * dgamma(lambda, ex_1_prior[[1]], ex_1_prior[[2]]) + (1 - p) * dgamma(lambda, ex_2_prior[[1]], ex_2_prior[[2]])
453 posterior <-
454   mcoefs[[1]] * dgamma(lambda, ex_1_prior[[1]] + n, ex_1_prior[[2]] + n *
455     mean_wait) + mcoefs[[2]] * dgamma(lambda, ex_2_prior[[1]] + n, ex_2_prior[[2]] + n *
456     mean_wait)
457
458 prlikpost <-
459   reshape2::melt(list(
460     Prior = prior,
461     Likelihood = likelihood,
462     Posterior = posterior
463   ))
464 ggcompplot <-
465   ggplot(data = prlikpost) + geom_area(
466     aes(
467       x = rep(lambda, 3),
468       y = value,
469       colour = L1,
470       fill = L1
471     ),
472     alpha = 0.3,
473     size = 0.5,
474     position = "identity"
475   ) +
476   theme_minimal() +
477   labs(x = "Lambda", y = "Density", title = "Prior-Posterior Compatibility Plot")
478 ggcompplot
479
480 mcoefs[[1]] *
481   Renext::qlomax(c(0.025, 0.975), betas[1], alphas[1]) +
482   mcoefs[[2]] *
483   Renext::qlomax(c(0.025, 0.975), betas[2], alphas[2])
484
485 dens_lind <- data.frame(postdraw = smps_lind$postdraw)
486 ggdens_lind <-
487   ggplot(data = dens, aes(x = postdraw)) +
488   geom_density(size = 1, color = "darkblue", fill = "lightblue") +
489   theme_minimal() +
490   labs(x = "Posterior Waiting Time", y = "Density of Posterior Predictive", title = "Posterior Predictive Density (Lindley Likelihood)") +
491   geom_vline(
492     xintercept = mean(dens_lind$postdraw),
493     size = 2,
494     color = "red"
495   ) +
496   geom_vline(
497     xintercept = quantile(dens_lind$postdraw, c(0.025, 0.975)),
498     color = "darkred",
499     size = 1
500   ) + coord_cartesian(xlim = c(0, 50))
501 ggdens_lind
502
503 quantile_set <- seq(0.01, 0.99, by = 0.01)
504 data_q <- quantile(waiting_times, quantile_set)
505 exp_q <- quantile(smps$postdraw, quantile_set)
506 lind_q <- quantile(smps_lind$postdraw, quantile_set)
507
508 quantile_df <- data.frame(obs = data_q, exp = exp_q, lind = lind_q)
509 ggqq <-
510   ggplot(data = quantile_df, aes(y = obs)) + coord_cartesian(xlim = c(0, 45), ylim = c(0, 45)) +
511   geom_point(aes(x = exp, colour = "Exponential")) +
512   geom_point(aes(x = lind, colour = "Lindley")) +
513   theme_minimal() +
514   labs(x = "Fitted Quantiles", y = "Sample Quantiles") + geom_abline(slope = 1, intercept = 0) + scale_colour_discrete(name = "Distribution")
515
516 ggqq

```

## C Stan Code for Question 2

### C.1 Exponential Likelihood

```

1 data {
2   // number of mixture components
3   int<lower=1> N; // number of data points
4   real y[N]; // observations
5   simplex[2] theta; // mixing proportions
6   positive_ordered[2] alpha; // locations of mixture components
7   vector<lower=0>[2] beta; // scales of mixture components
8 }
9 parameters {
10   real<lower=0> lambda;
11 }
12 model {
13   y ~ exponential(lambda);
14   target += log_mix(theta[1],
15     gamma_lpdf(lambda | alpha[1], beta[1]),
16     gamma_lpdf(lambda | alpha[2], beta[2]));
17 }
18 generated quantities {
19   real<lower=0> mwt = 1 / lambda;
20   real postdraw = exponential_rng(lambda);

```

```
21 }
```

## C.2 Lindley Likelihood

```
1 functions{
2   real lindley_lpdf(real y, real lambda){
3     real lpdf = 2*log(lambda) - log(lambda+1) + log(1+y) - lambda * y;
4     return lpdf;
5   }
6   real lindley_rng(real lambda){
7     real u = uniform_rng(0,1);
8     real p = lambda/(lambda+1);
9     if (u<=p){
10      real v = exponential_rng(lambda);
11      return v;
12    }
13    else{
14      real w = gamma_rng(2, lambda);
15      return w;
16    }
17  }
18 }
19 data {
20   // number of mixture components
21   int<lower=1> N; // number of data points
22   real y[N]; // observations
23   simplex[2] theta; // mixing proportions
24   positive_ordered[2] alpha; // locations of mixture components
25   vector<lower=0>[2] beta; // scales of mixture components
26 }
27 parameters {
28   real<lower=0> lambda;
29 }
30 model {
31   for (n in 1:N){
32     y[n] ~ lindley(lambda);
33   }
34   target += log_mix(theta[1],
35     gamma_lpdf(lambda | alpha[1], beta[1]),
36     gamma_lpdf(lambda | alpha[2], beta[2]));
37 }
38 generated quantities {
39   real<lower=0> mvt = ((lambda + 2)/(lambda*(lambda + 1)));
40   real postdraw = lindley_rng(lambda);
41 }
```

## D JAGS Code for Question 2

```
1 model{
2   for (i in 1:n) {
3     y[i] ~ dexp(lambda)
4   }
5   lambda ~ dgamma(a[pick], b[pick]) # a[1]=9.2, b[1]=13.8; a[2]=12; b[2] = 3
6   pick ~ dcat(p[1:2]) # pick takes value 1 or 2 with prior prob p[1] or p[2]
7   ypred ~ dexp(lambda)
8   texp = 1 / lambda
9 }
```

## E R Code for Question 3

```
1 library(coda)
2 library(rstan)
3 library(rstantools)
4 library(bayesplot)
5 library(ggplot2)
6
7 library(dplyr)
8 library(data.table)
9 library(modeest)
10
11 rstan_options(auto_write = TRUE)
12 options(datatable fread.datatable = FALSE)
13 options(mc.cores = parallel::detectCores() - 1)
14 devAskNewPage(ask = FALSE)
15 par(ask = F)
16 #options(mc.cores = 1)
17
18 par(mfrow = c(1,1))
19
20 abalone <- fread('data/abalone.csv')
21 names(abalone) <- make.names(names(abalone))
22 abalone$Age <- abalone$Rings + 1.5
23 abalone <- subset(abalone, select = -c(Rings))
24
25 cmeans <-
26   colMeans(Filter(is.numeric, subset(abalone, select = -c(Age))))
27 amean <- mean(abalone$Age)
28
```



```

29 scale2 <- function(x, na.rm = FALSE)
30   (x - mean(x, na.rm = na.rm))
31 abalone <- mutate_if(abalone, is.numeric, scale2, na.rm = TRUE)
32 abalone$Sex <- as.factor(abalone$Sex)
33
34 X <-
35   model.matrix(
36     Age ~ Sex + Length + Diameter + Height + Whole.weight + Shucked.weight + Viscera.weight + Shell.weight,
37     data = abalone
38   )
39
40 xlevs <-
41   lapply(abalone[, sapply(abalone, is.factor), drop = F], function(j) {
42     levels(j)
43   })
44
45 new_obs <-
46   data.frame(
47     Sex = 'M',
48     Length = 0.515,
49     Diameter = 0.400,
50     Height = 0.133,
51     Whole.weight = 0.531,
52     Shucked.weight = 0.231,
53     Viscera.weight = 0.122,
54     Shell.weight = 0.168
55   )
56
57 dfcmean <- as.data.frame(t(cmeans))
58
59 new_obs[-1] <- new_obs[-1] - dfcmean
60 mm_new <- model.matrix(~ ., data = new_obs, xlev = xlevs)
61
62 f_lm <- lm(data = abalone, Age ~ .)
63 summary(f_lm)
64 pred_f_lm <- predict(f_lm, new_obs)
65 pred_f_lm
66 par(mfrow = c(2, 2))
67 plot(f_lm)
68 par(mfrow = c(1, 1))
69 plot(fitted(f_lm), abalone$Age, asp = 1)
70
71
72 b_lm_data <-
73   list(
74     N = nrow(X),
75     K = ncol(X),
76     X = X,
77     y = abalone$Age,
78     N_new = 1,
79     x_new = as.array(mm_new),
80     p_params = c(1e4, 1e-1, 1e-1, 1/27)
81   )
82
83 b_lm <- stan_model(file = 'aiq3.stan')
84
85 b_lm_fit <- sampling(
86   b_lm,
87   data = b_lm_data,
88   chains = 7,
89   control = list(adapt_delta = 0.8),
90   iter = 16000
91 )
92 plot(b_lm_fit)
93 b_lm_fit
94
95 smps <- extract(b_lm_fit)
96 # tp <- traceplot(b_lm_fit, pars = c("beta", "sigma"))
97 # tp
98
99 b_lm_coda <- As.mcmc.list(b_lm_fit)
100 gelman.plot(b_lm_coda, ask=FALSE)
101 # gelman.diag(b_lm_coda)
102 # plot(b_lm_coda, ask = FALSE)
103
104 eap_beta <- matrix(colMeans(smps$beta), ncol = 1)
105
106 b_lm_age <- X %>% eap_beta
107
108 n_dens <-
109   ggplot(data = data.frame(samp = smps$y_new, aes(x = samp))) + geom_density() + coord_cartesian(xlim = c(-10,10))
110 n_dens
111
112 pred_eap <- mean(smps$y_new)
113 pred_map <- mlv(smps$y_new, method = "venter")
114 pred_f_lm
115 pred_map
116 pred_eap
117
118 available_mcmc(pattern = "_nuts_")
119
120 posterior_lm <- as.array(b_lm_fit)
121 np_lm <- nuts_params(b_lm_fit)
122 lp_lm <- log_posterior(b_lm_fit)
123
124 mcmc_nuts_divergence(np_lm, lp_lm)
125 mcmc_nuts_acceptance(np_lm, lp_lm)
126
127 # pairs(b_lm_fit, pars = c("beta[1]", "beta[2]", "beta[3]"))
128 # pairs(b_lm_fit, pars = c("beta[4]", "beta[5]", "beta[6]"))
129 # pairs(b_lm_fit, pars = c("beta[7]", "beta[8]", "beta[9]", "beta[10]"))
130
131 color_scheme_set("viridisD")
132 mcmc_trace(posterior_lm, pars = c("beta[1]", "beta[2]", "beta[3]", "beta[4]", "beta[5]", "beta[6]", "beta[7]", "beta[8]", "beta[9]", "beta[10]", "sigma"), np = np_lm, size = 0.05)
133
134 # b_lm_data_test <-

```

```

135 # list(
136 #   N = nrow(X),
137 #   K = ncol(X),
138 #   X = X,
139 #   y = abalone$Age,
140 #   N_new = 1,
141 #   x_new = as.array(mm_new),
142 #   p_params = c(1e1, 7, 30)
143 # )
144 #
145 # b_lm_fit_prior_sensitivity <- sampling(
146 #   b_lm,
147 #   data = b_lm_data_test,
148 #   chains = 7,
149 #   control = list(adapt_delta = 0.8),
150 #   iter = 3000
151 # )
152 #
153 # b_lm_fit_prior_sensitivity
154 #
155 # posterior_lm_psa <- as.array(b_lm_fit_prior_sensitivity)
156 # np_lm_psa <- nuts_params(b_lm_fit_prior_sensitivity)
157 # lp_lm_psa <- log_posterior(b_lm_fit_prior_sensitivity)
158 #
159 # mcmc_nuts_divergence(np_lm_psa, lp_lm_psa)
160 # mcmc_nuts_acceptance(np_lm_psa, lp_lm_psa)
161 #
162 # pairs(b_lm_fit_prior_sensitivity, pars = c("beta[1]", "beta[2]", "beta[3]"))
163 # pairs(b_lm_fit_prior_sensitivity, pars = c("beta[4]", "beta[5]", "beta[6]"))
164 # pairs(b_lm_fit_prior_sensitivity, pars = c("beta[7]", "beta[8]", "beta[9]", "beta[10]"))
165 #
166 # color_scheme_set("viridisD")
167 # mcmc_trace(posterior_lm_psa, pars = c("beta[1]", "beta[2]", "beta[3]", "beta[4]", "beta[5]", "beta[6]", "beta[7]", "beta[8]", "beta[9]", "beta
168 #   [10]", "sigma"), np = np_lm_psa, size = 0.05)
169 #
170 quantile_set <- seq(0.01, 0.99, by = 0.01)
171 data_q <- quantile(abalone$Age, quantile_set)
172 bayes_q <- quantile(b_lm_age, quantile_set)
173 lm_q <- quantile(fitted(f_lm), quantile_set)
174 #
175 quantile_df <- data.frame(obs = data_q, b = bayes_q, l = lm_q)
176 ggqq <- ggplot(data = quantile_df, aes(y = obs)) + coord_cartesian(xlim = c(-10, 10), ylim = c(-10, 10)) +
177   geom_point(aes(x = b, colour = "Bayesian Robust Linear Model")) +
178   geom_point(aes(x = l, colour = "Normal Linear Model")) +
179   theme_minimal() +
180   labs(x = "Fitted Quantiles", y = "Sample Quantiles") + geom_abline(slope = 1, intercept = 0) + scale_colour_discrete(name = "Distribution")
181 ggqq
182 #####
183 #####
184 #
185 # b_lm_normal <- stan_model(file = 'aiq3normal.stan')
186 #
187 # b_lm_fit_normal <- sampling(
188 #   b_lm_normal,
189 #   data = b_lm_data,
190 #   chains = 7,
191 #   control = list(adapt_delta = 0.8),
192 #   iter = 4000
193 # )
194 # plot(b_lm_fit_normal)
195 # b_lm_fit_normal
196 #
197 # smps_norm <- extract(b_lm_fit_normal)
198 #
199 # eap_beta_norm <- matrix(colMeans(smps_norm$beta), ncol = 1)
200 #
201 # b_lm_age_norm <- X %*% eap_beta_norm
202 #
203 # n_dens_norm <-
204 #   ggplot(data = data.frame(samp = smps_norm$y_new), aes(x = samp)) + geom_density() + coord_cartesian(xlim = c(-10,10))
205 # n_dens_norm
206 #
207 # quantile_set <- seq(0.01, 0.99, by = 0.01)
208 # data_q <- quantile(abalone$Age, quantile_set)
209 # bayes_q <- quantile(b_lm_age, quantile_set)
210 # lm_q <- quantile(fitted(f_lm), quantile_set)
211 # bayes_n_q <- quantile(b_lm_age_norm, quantile_set)
212 #
213 # quantile_df <- data.frame(obs = data_q, b = bayes_q, l = lm_q, bn = bayes_n_q)
214 # ggqq <- ggplot(data = quantile_df, aes(y = obs)) + coord_cartesian(xlim = c(-10, 10), ylim = c(-10, 10)) +
215 #   geom_point(aes(x = b, colour = "Bayesian Robust Linear Model")) +
216 #   geom_point(aes(x = l, colour = "Normal Linear Model")) +
217 #   geom_point(aes(x = bn, colour = "Bayesian Normal Linear Model")) +
218 #   theme_minimal() +
219 #   labs(x = "Fitted Quantiles", y = "Sample Quantiles") + geom_abline(slope = 1, intercept = 0) + scale_colour_discrete(name = "Distribution")
220 # ggqq
221 #
222 # n_dens <-
223 #   ggplot(data = data.frame(samp = smps$y_new), aes(x = samp)) +
224 #   geom_density(size = 1, color = "darkblue", fill = "lightblue", alpha = 0.3) +
225 #   coord_cartesian(xlim = c(-10,10)) +
226 #   theme_minimal() +
227 #   labs(title = "Posterior Predictive for New Data", x = "Mean-Corrected Age", y = "Density")
228 # n_dens

```

## F Stan Code for Question 3

### F.1 Linear Model with $t$ Errors

```

1 data {
2   int<lower=1> K;
3   int<lower=0> N;
4   matrix[N, K] X;
5   vector[N] y;
6   int<lower=0> N_new;
7   matrix[N_new, K] x_new;
8   vector[4] p_params;
9 }
10 parameters {
11   vector[K] beta;
12   real<lower=0> sigma;
13   real<lower=0> nu;
14 }
15 model {
16   nu ~ exponential(p_params[4]);
17   beta ~ normal(0, p_params[1]);
18   sigma ~ inv_gamma(p_params[2], p_params[3]);
19   y ~ student_t(nu+1, X * beta, sigma);
20 }
21 generated quantities {
22   vector[N_new] y_new;
23   for (n in 1:N_new)
24     y_new[n] = student_t_rng(nu+1, x_new[n] * beta, sigma);
25 }

```

## F.2 Linear Model with Normal Errors

```

1 data {
2   int<lower=1> K;
3   int<lower=0> N;
4   matrix[N, K] X;
5   vector[N] y;
6   int<lower=0> N_new;
7   matrix[N_new, K] x_new;
8   vector[3] p_params;
9 }
10 parameters {
11   vector[K] beta;
12   real<lower=0> sigma;
13 }
14 model {
15   beta ~ normal(0, p_params[1]);
16   sigma ~ inv_gamma(p_params[2], p_params[3]);
17   y ~ normal(X * beta, sigma);
18 }
19 generated quantities {
20   vector[N_new] y_new;
21   for (n in 1:N_new)
22     y_new[n] = normal_rng(x_new[n] * beta, sigma);
23 }

```